


Implemented IntrMotiv families: signal and learning paths

Common backbone · R = 8, L = 64 · no learned recurrent matrix



Solid: information or recorded feedback Dashed green: optimization path

CPD contextual variant
Past CA3 ($\pm$ actions) gates DG logits
DIRECT detaches; BPTT retains history
Temporal candidate: no predictor

W_REF variant
Separate frozen SCR detector defines arrival
Live DG continues intrinsic learning
STOP / JOINT changes PPO → DG only

Family differences

SCR / SAT: ARR credit + C15 graph; PPO stops at DG

DGP: HIT / FIRST × STOP / JOINT × LEG / FiLM

PIC F_GATE: no goal, no graph · PIC G_SHARED / W_REF: goal, no graph